

Neural Net Implementation on FPGA's: Goals

Cory Nezin, Brenda So

September 6, 2017

1. Design and implement hardware optimized inference on an FPGA.
 - (a) Create a framework to convert symbolic neural nets to FPGA code.
 - i. ReLU, Sigmoid, Threshold.
 - ii. Fully Connected (matrix multiply - with Strassen algorithm?).
 - iii. Sparsely Connected? (Sparse matrix multiply).
 - iv. 1D (2D and 3D also?) Convolution - just FIR filters.
 - v. Max pooling, average pooling, downsampling.
 - vi. Recurrent elements?
 - (b) Implement weight compression to decrease neural net size.
 - i. Use multiplierless multiplication?
 - ii. Use SVD for fully connected compression?
 - iii. Implement Optimal Brain Damage.
 - iv. Implement Optimal Brain Surgeon.
 - v. Implement Deep Compression for specialized hardware.
 - (c) Test neural net on a toy data set (spiral, checkerboard, wdbc).
 - i. Implement a simple live input method - ADC?
 - ii. Perform live classification of hand written digits.
2. Benchmark performance of FPGA's against GPU's and CPU's.
 - (a) Quantify operation in floating point and fixed point.
 - (b) Measure throughput, power, accuracy, memory, and latency.

3. Build a general front end, pick one of the following:
- (a) Radar - Gender identification (toy problem), people in room.
 - (b) Style Transfer - live feed.
 - (c) Audio - Frank's speech denoising.
 - (d) Video - Live lip reading, live blackboard transcription.
 - (e) Image - Handwriting transcription.
 - (f) Lidar - Car or pedestrian detection.
 - (g) Wireless - Signal classification, smart jamming?